

```
In [162... # Data manipulation and visualization
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path
import os
import warnings
warnings.filterwarnings('ignore')

# Deep Learning Libraries
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, models
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D, Input
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img, img_to_array
from tensorflow.keras.applications import VGG16, ResNet50, MobileNetV2
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau

# Scikit-Learn for metrics
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.metrics import ConfusionMatrixDisplay

# Set random seeds for reproducibility
np.random.seed(42)
tf.random.set_seed(42)
```

```
In [64]: patient_df = pd.read_csv('./skin-cancer-mnist-ham10000/HAM10000_metadata.csv')
```

Checking the Data & Cleaning

```
In [65]: patient_df
```

Out[65]:

	lesion_id	image_id	dx	dx_type	age	sex	localization
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear
...	...	...	...	...	...	...	...
10010	HAM_0002867	ISIC_0033084	akiec	histo	40.0	male	abdomen
10011	HAM_0002867	ISIC_0033550	akiec	histo	40.0	male	abdomen
10012	HAM_0002867	ISIC_0033536	akiec	histo	40.0	male	abdomen
10013	HAM_0000239	ISIC_0032854	akiec	histo	80.0	male	face
10014	HAM_0003521	ISIC_0032258	mel	histo	70.0	female	back

10015 rows × 7 columns

In [66]: `patient_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10015 entries, 0 to 10014
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   lesion_id       10015 non-null  object
1   image_id        10015 non-null  object
2   dx              10015 non-null  object
3   dx_type         10015 non-null  object
4   age            9958 non-null   float64
5   sex            10015 non-null  object
6   localization    10015 non-null  object
dtypes: float64(1), object(6)
memory usage: 547.8+ KB
```

In [67]: `patient_df.describe()`

Out[67]:

	age
count	9958.000000
mean	51.863828
std	16.968614
min	0.000000
25%	40.000000
50%	50.000000
75%	65.000000
max	85.000000

Based on these summaries of the dataset, there are 10015 rows of data within the HAM10000 skin lesion dataset.

The dataset contains information on the lesion and image id, the type of disease the patient has (dx column), the type of test the patient had (dx_type column), the patient's age, sex and where the skin lesion is located on.

Age is the only numerical variable in the entire dataset, and the mean age is roughly 51.8 years old.

```
In [68]: patient_df.isna().sum()
```

```
Out[68]: lesion_id      0
image_id      0
dx            0
dx_type      0
age          57
sex          0
localization  0
dtype: int64
```

```
In [69]: patient_df.duplicated().sum()
```

```
Out[69]: 0
```

```
In [70]: patient_df['age'] = patient_df['age'].fillna(patient_df['age'].mean().round(1))
```

There are 57 missing values for the age column of the dataset, so I filled up all the missing values with the mean of the overall patients' age.

```
In [71]: patient_df.isna().sum()
```

```
Out[71]: lesion_id      0
image_id      0
dx            0
dx_type      0
age           0
sex           0
localization  0
dtype: int64
```

Below I mapped out readable labels to better identify all the different skin lesions.

```
In [72]: lesion_type_dict = {
        'nv': 'Melanocytic nevi',
        'mel': 'Melanoma',
        'bkl': 'Benign keratosis-like lesions ',
        'bcc': 'Basal cell carcinoma',
        'akiec': 'Actinic keratoses',
        'vasc': 'Vascular lesions',
        'df': 'Dermatofibroma'
    }
```

```
In [73]: patient_df['cell_type'] = patient_df['dx'].map(lesion_type_dict)
patient_df.head()
```

Out[73]:

	lesion_id	image_id	dx	dx_type	age	sex	localization	cell_type
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp	Benign keratosis-like lesions
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp	Benign keratosis-like lesions
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp	Benign keratosis-like lesions
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp	Benign keratosis-like lesions
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear	Benign keratosis-like lesions

Below shows a few invalid records (the patient's age is 0 and their gender is unknown), I filtered and removed these records from the dataset to ensure data integrity.

In [74]: `patient_df['age'].value_counts()`

Out[74]:

```

age
45.0    1299
50.0    1187
55.0    1009
40.0     985
60.0     803
70.0     756
35.0     753
65.0     731
75.0     618
30.0     464
80.0     404
85.0     290
25.0     247
20.0     169
5.0       86
15.0      77
51.9      57
10.0      41
0.0       39
Name: count, dtype: int64

```

In [75]: `patient_df = patient_df[patient_df['age']!=0]`

In [76]: `patient_df['sex'].value_counts()`

Out[76]:

```

sex
male      5383
female    4538
unknown     55
Name: count, dtype: int64

```

```
In [77]: patient_df = patient_df[patient_df['sex']!='unknown']
```

```
In [196... patient_df.head()
```


Out[196...

	lesion_id	image_id	dx	dx_type	age	sex	localization	cell_type	im
0	HAM_0000118	ISIC_0027419	bkl	histo	80.0	male	scalp	Benign keratosis- like lesions	ISIC_00:
1	HAM_0000118	ISIC_0025030	bkl	histo	80.0	male	scalp	Benign keratosis- like lesions	ISIC_00:
2	HAM_0002730	ISIC_0026769	bkl	histo	80.0	male	scalp	Benign keratosis- like lesions	ISIC_00:
3	HAM_0002730	ISIC_0025661	bkl	histo	80.0	male	scalp	Benign keratosis- like lesions	ISIC_00:
4	HAM_0001466	ISIC_0031633	bkl	histo	75.0	male	ear	Benign keratosis- like lesions	ISIC_00:

◀ ▶

Exploratory Data Analysis (EDA)

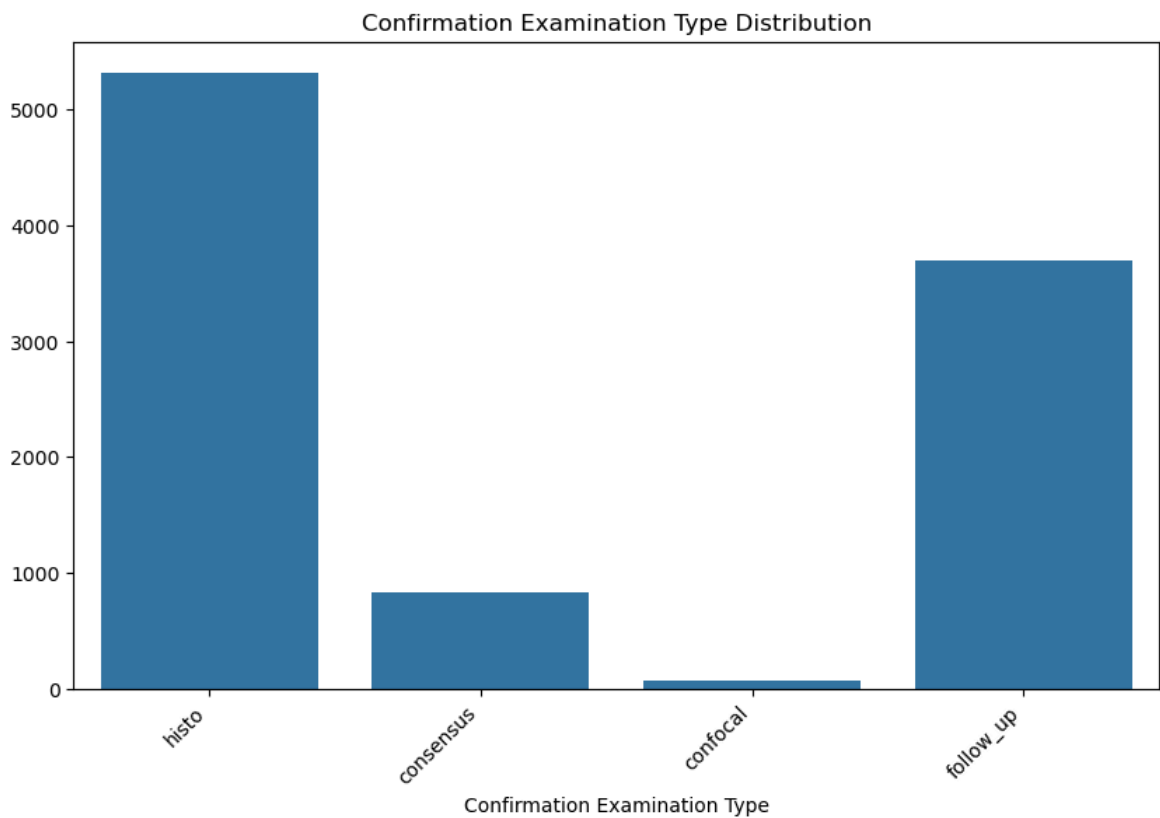
```
In [103... plt.figure(figsize=(10,6))
sns.countplot(data=patient_df, x='cell_type')
plt.xticks(rotation=45,ha='right')
plt.title('Skin Lesion Type Distribution')
plt.xlabel('Skin Lesion Type')
plt.ylabel('')
plt.show()
```



A majority of patients have Melanocytic Nevi, a benign skin lesion.

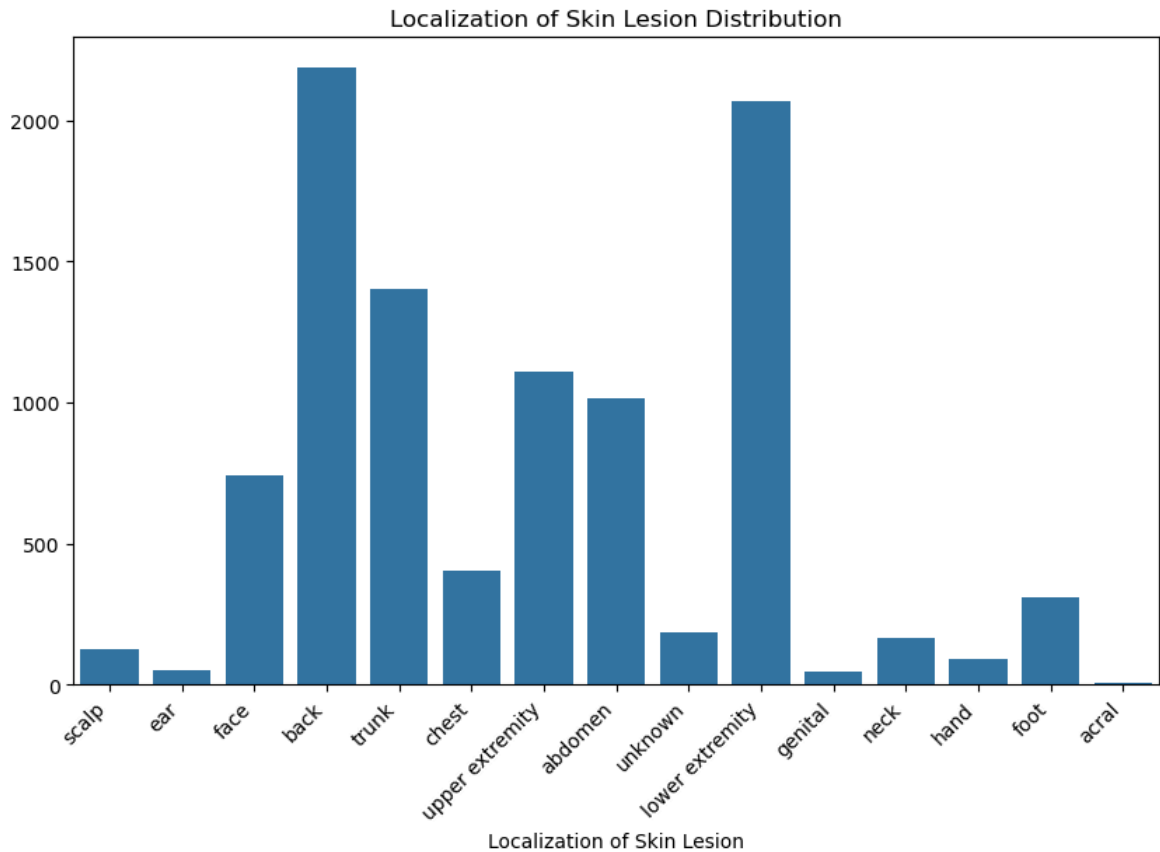
The type of skin lesion that is least identified in patients is Dermatofibroma, also a benign skin lesion.

```
In [102... plt.figure(figsize=(10,6))
sns.countplot(data=patient_df, x='dx_type')
plt.xticks(rotation=45,ha='right')
plt.title('Confirmation Examination Type Distribution')
plt.xlabel('Confirmation Examination Type')
plt.ylabel('')
plt.show()
```



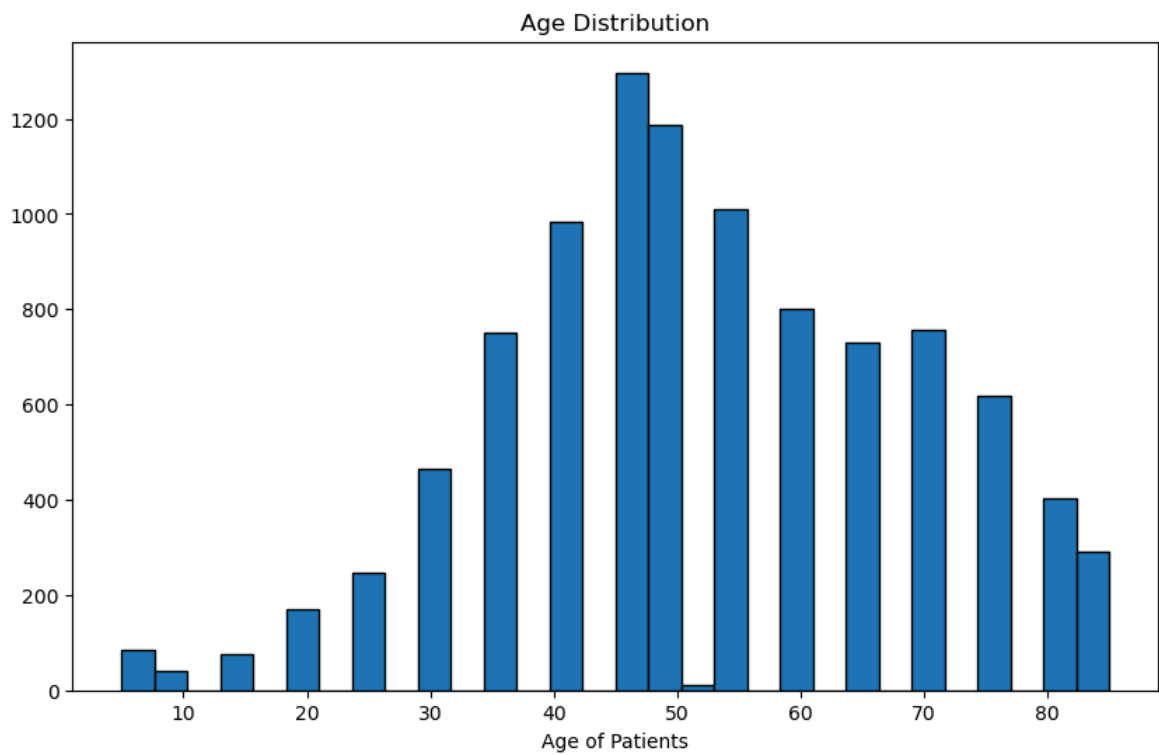
Most patients confirmed the type of skin lesion they had through a histopathology (histo), whereas the least amount of patients got confirmation by in-vivo confocal microscopy (confocal)

```
In [101... plt.figure(figsize=(10,6))
sns.countplot(data=patient_df, x='localization')
plt.xticks(rotation=45,ha='right')
plt.title('Localization of Skin Lesion Distribution')
plt.xlabel('Localization of Skin Lesion')
plt.ylabel('')
plt.show()
```



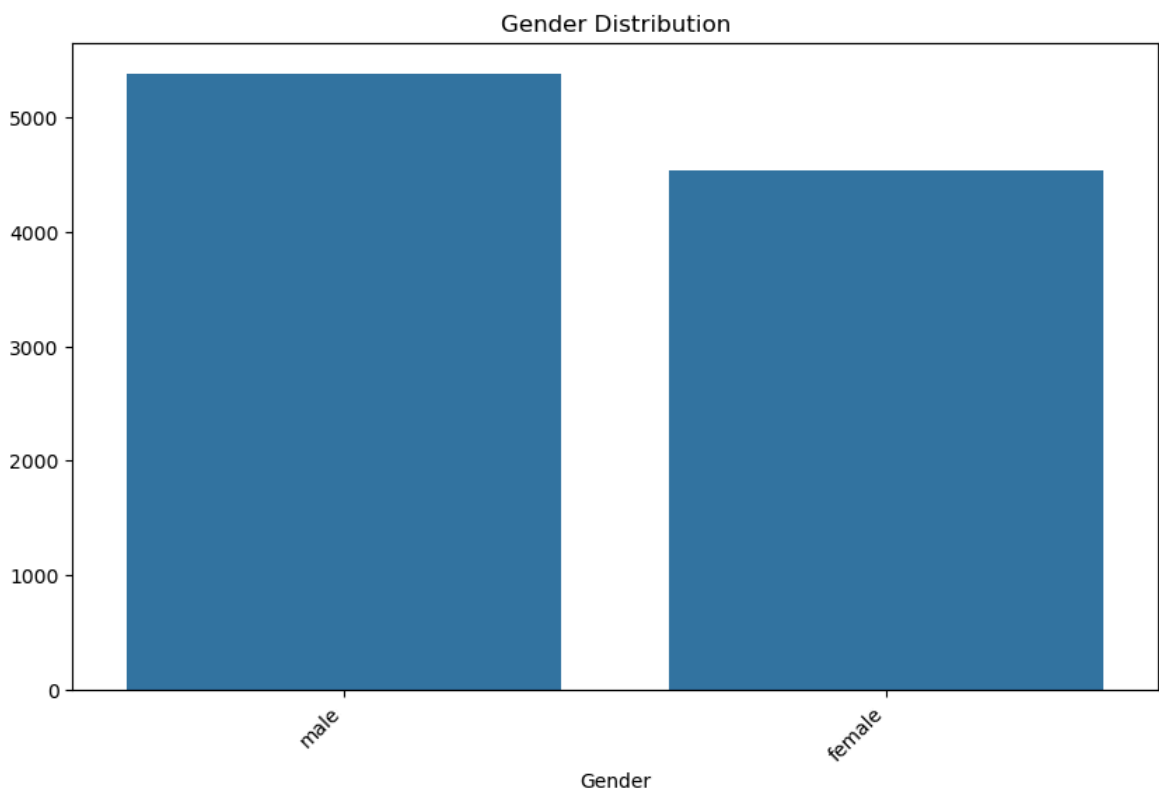
Most skin lesions tend to grow on the back, whereas only 7 patients have been confirmed to have skin lesions on their acral body parts, which are the hands, feet, ears or nose of a patient.

```
In [104... plt.figure(figsize=(10,6))
plt.hist(x=patient_df['age'], bins=30, edgecolor='black')
plt.title('Age Distribution')
plt.xlabel('Age of Patients')
plt.ylabel('')
plt.show()
```

Most patients are around the age of 35 to 60 years old.

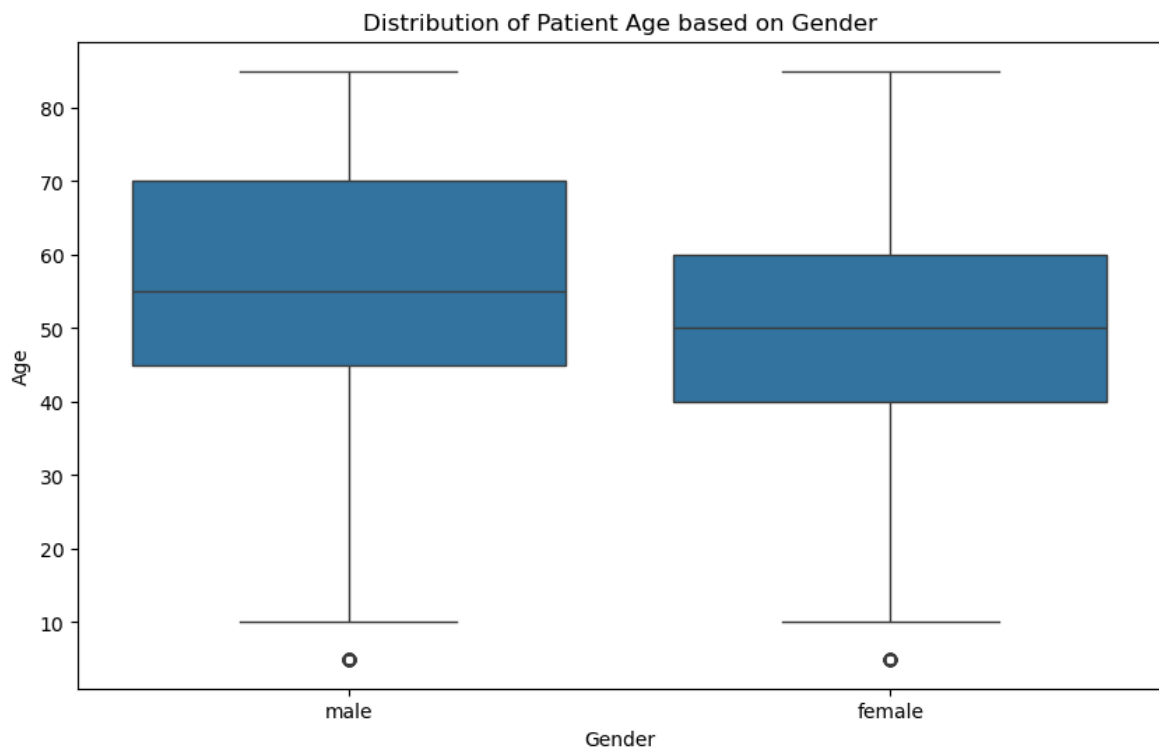
```
In [105... plt.figure(figsize=(10,6))
sns.countplot(data=patient_df, x='sex')
plt.xticks(rotation=45,ha='right')
plt.title('Gender Distribution')
plt.xlabel('Gender')
plt.ylabel('')
plt.show()
```



Most patients are male.

In [121...

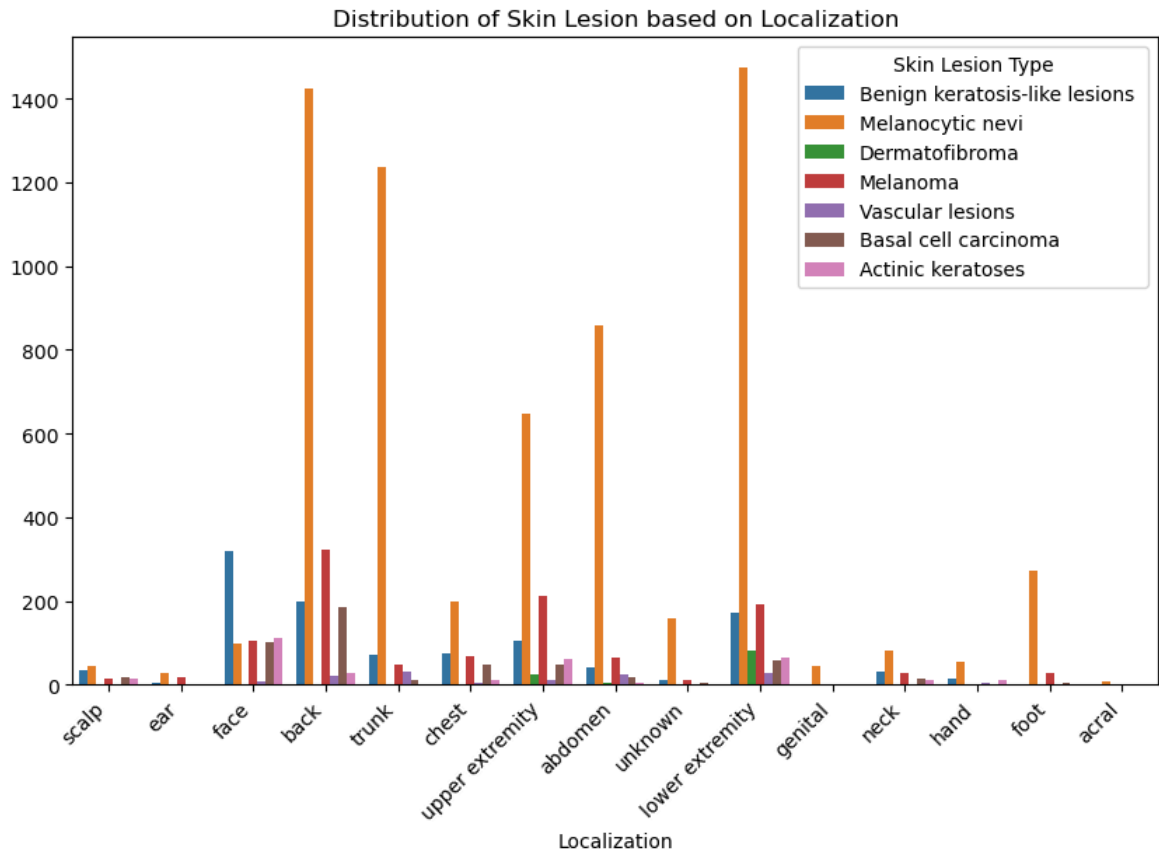
```
plt.figure(figsize=(10,6))
sns.boxplot(data=patient_df, x='sex', y='age')
plt.title('Distribution of Patient Age based on Gender')
plt.xlabel('Gender')
plt.ylabel('Age')
plt.show()
```



Most male patients are between the age group of 45 - 70 years old while most female patients are between the ages of 40 - 60 years old.

In [98]:

```
plt.figure(figsize=(10,6))
sns.countplot(data=patient_df, x='localization', hue='cell_type')
plt.xticks(rotation=45, ha='right')
plt.title('Distribution of Skin Lesion based on Localization')
plt.xlabel('Localization')
plt.ylabel('')
plt.legend(title='Skin Lesion Type', bbox_to_anchor=(1,1))
plt.show()
```

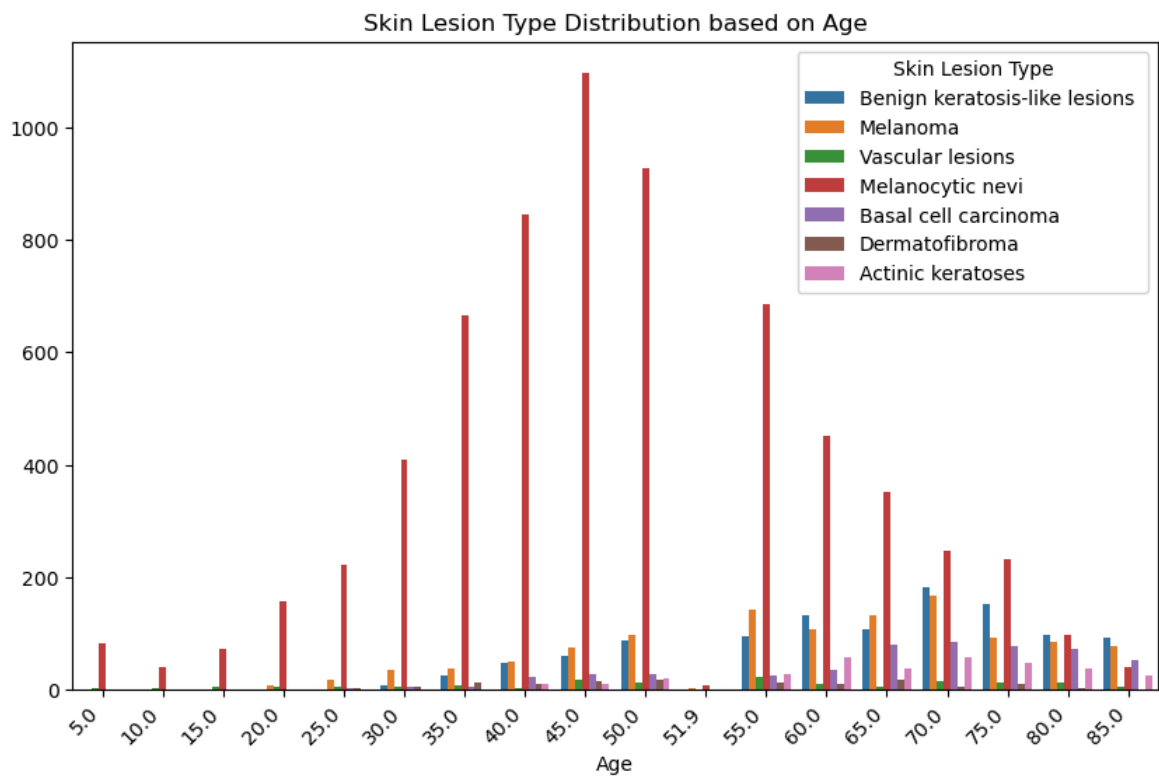


Melanocytic nevi is most commonly diagnosed in all areas of the body except for the face.

The majority of skin lesions found on the face are Benign keratosis-like lesions.

In [109...

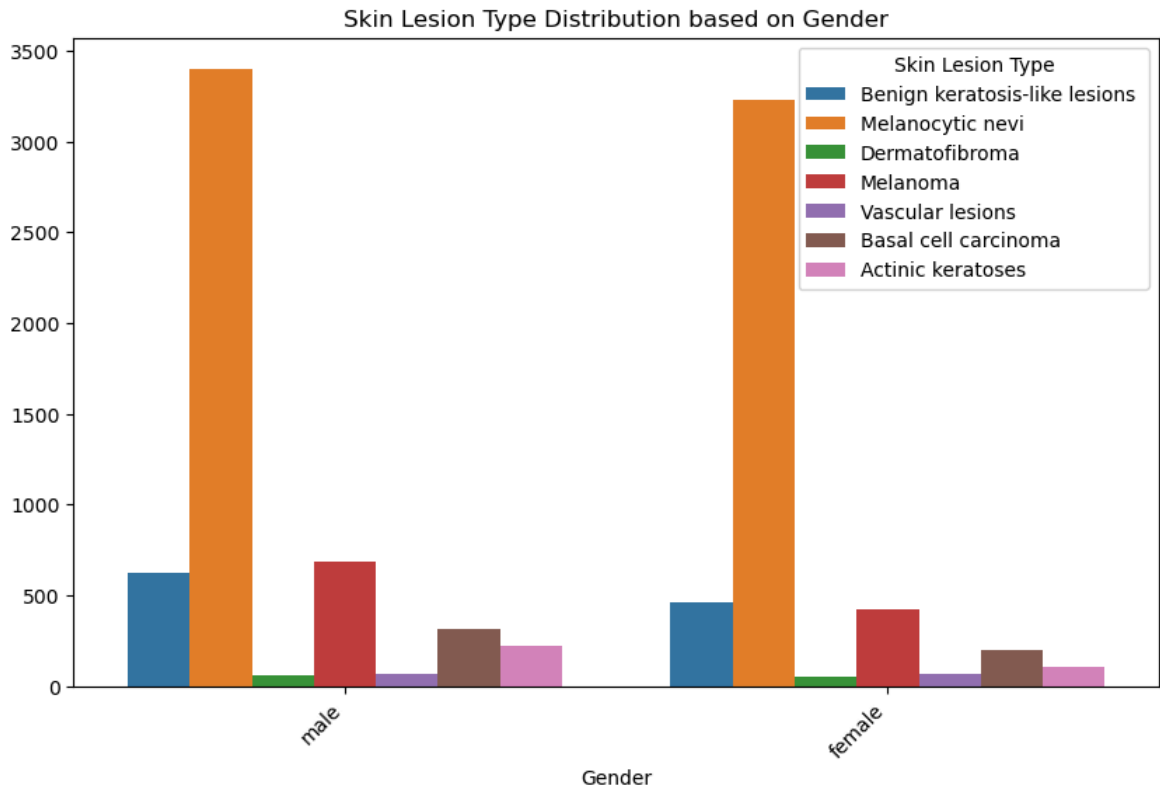
```
plt.figure(figsize=(10,6))
sns.countplot(data=patient_df, x='age', hue='cell_type')
plt.xticks(rotation=45,ha='right')
plt.title('Skin Lesion Type Distribution based on Age')
plt.xlabel('Age')
plt.ylabel('')
plt.legend(title='Skin Lesion Type',bbox_to_anchor=(1,1))
plt.show()
```



Melanocytic nevi is most commonly diagnosed across all age groups except for 80 years old and above.

For patients within the age groups of 80 and above, Benign keratosis-like lesions are more common.

```
In [110... plt.figure(figsize=(10,6))
sns.countplot(data=patient_df, x='sex', hue='cell_type')
plt.xticks(rotation=45,ha='right')
plt.title('Skin Lesion Type Distribution based on Gender')
plt.xlabel('Gender')
plt.ylabel('')
plt.legend(title='Skin Lesion Type',bbox_to_anchor=(1,1))
plt.show()
```



Melanocytic nevi is most commonly diagnosed among both male and female patients.

Preparing the Data

Creating an image path to link the image with its corresponding data in the dataframe

```
In [112...] patient_df['image_path'] = patient_df['image_id'].astype(str) + '.jpg'
```

Preparing the directory paths and setting up image parameters for processing

```
In [113...] BASE_DIR = Path('./skin-cancer-mnist-ham10000') # Update this to your data dire
TRAIN_DIR = BASE_DIR / 'train'
TEST_DIR = BASE_DIR / 'test'
SAMPLE_DIR = BASE_DIR / 'sample'

# Image parameters
IMG_HEIGHT = 224
IMG_WIDTH = 224
IMG_CHANNELS = 3
BATCH_SIZE = 32
TARGET_SIZE = (IMG_HEIGHT, IMG_WIDTH)

print(f"{TRAIN_DIR}\nTrain directory exists: {TRAIN_DIR.exists()}")
print(f"{TEST_DIR}\nTest directory exists: {TEST_DIR.exists()}")
print(f"{SAMPLE_DIR}\nSample directory exists: {SAMPLE_DIR.exists()}")
```

```
skin-cancer-mnist-ham10000\train
Train directory exists: True
skin-cancer-mnist-ham10000\test
Test directory exists: True
skin-cancer-mnist-ham10000\sample
Sample directory exists: True
```

```
In [114... # Count images in each set
def count_images(directory):
    return len(list(directory.glob('*.jpg')))

# Count images in each set
train_total = count_images(TRAIN_DIR)
test_total = count_images(TEST_DIR)

print(f"Training images: {train_total}")
print(f"Test images: {test_total}")
```

Training images: 8000

Test images: 2015

Sample a few images from the training directory.

```
In [115... def plot_sample_images(directory, num_samples=12, figsize=(10, 6)):
    fig, axes = plt.subplots(2, 3, figsize=figsize)
    axes = axes.ravel()
    images = list(directory.glob('*.jpg'))


    for idx in range(0,6):
        if images:
            random_image = np.random.choice(images)
            img = load_img(random_image, target_size=(IMG_HEIGHT, IMG_WIDTH))

            axes[idx].imshow(img)
            axes[idx].set_title(random_image.stem, fontsize=10)
            axes[idx].axis('off')

    plt.tight_layout()
    plt.suptitle('Sample Skin Cancer Images', fontsize=16, y=1.02)
    plt.show()

plot_sample_images(TRAIN_DIR)
```

Sample Skin Cancer Images



Sample Skin Cancer Images with Class

```
In [192... def plot_sample_images_by_class(directory, num_samples=12, figsize=(10, 6)):
    fig, axes = plt.subplots(2, 3, figsize=figsize)
    axes = axes.ravel()
    images = list(directory.glob('*.jpg'))

    for idx in range(0,6):
        if images:
            random_image = np.random.choice(images)
            img = load_img(random_image, target_size=(IMG_HEIGHT, IMG_WIDTH))

            label = patient_df[patient_df['image_id']==random_image.stem]

            label_text = label['cell_type'].values[0]

            axes[idx].imshow(img)
            axes[idx].set_title(label_text, fontsize=10)
            axes[idx].axis('off')

    plt.tight_layout()
    plt.suptitle('Sample Skin Cancer Images', fontsize=16, y=1.02)
    plt.show()

plot_sample_images_by_class(TRAIN_DIR)
```



Data Augmentation

```
In [153... train_data = ImageDataGenerator(
    rescale=1./255, # normalize pixels values to 0-1
    rotation_range=40, # rotating images up 40 degrees
    width_shift_range=0.2, # randomly shifting horizontally by 20%
```

```

height_shift_range=0.2, # randomly shifting vertically by 20%
shear_range=0.2, # shear transformation
zoom_range=0.2, # random zoom by 20%
horizontal_flip=True, # random horizontal flip
)

test_data = ImageDataGenerator(
    rescale=1./255 # normalize pixels values to 0-1
)

```

Training & Testing Generators for the Model

```

In [143... # AI GENERATED code for getting data of images from the dataframe
train_gen = train_data.flow_from_dataframe(
    dataframe=patient_df,
    directory=TRAIN_DIR,
    x_col='image_path',
    y_col='cell_type',
    target_size=TARGET_SIZE,
    class_mode='sparse',
    batch_size=BATCH_SIZE
)

test_gen = test_data.flow_from_dataframe(
    dataframe=patient_df,
    directory=TEST_DIR,
    x_col='image_path',
    y_col='cell_type',
    target_size=TARGET_SIZE,
    class_mode='sparse',
    batch_size=BATCH_SIZE
)

## get class labels and number of classes from generators
class_names = list(test_gen.class_indices.keys())
num_classes = len(class_names)

```

Found 7929 validated image filenames belonging to 7 classes.
Found 1992 validated image filenames belonging to 7 classes.

Visualizing the Data Augmentation

```

In [151... def visualize_augmentation(generator, num_images=5):
    """Visualize augmented versions of a single image."""
    # Get a batch of images
    x_batch, y_batch = next(generator)

    # Take the first image
    img = x_batch[0]

    # Create augmented versions
    fig, axes = plt.subplots(1, num_images, figsize=(15, 3))

    axes[0].imshow(img)
    axes[0].set_title('Original')
    axes[0].axis('off')

```



```

# Generate augmented versions
img_array = img * 255.0 ## image has already been rescaled in train_data, so
img_array = img_array.reshape((1,) + img_array.shape)


aug_iter = train_data.flow(img_array, batch_size=1)

for i in range(1, num_images):
    batch = next(aug_iter)
    aug_img = batch[0]
    axes[i].imshow(aug_img)
    axes[i].set_title(f'Augmented {i}')
    axes[i].axis('off')

plt.tight_layout()
plt.show()

visualize_augmentation(train_gen)

```



Building a simple CNN model

```

In [172... simple_model = Sequential([
    Conv2D(filters=32, kernel_size=(3,3), activation='relu', input_shape=(224, 2
    MaxPooling2D(pool_size=(2,2)),
    Conv2D(filters=64, kernel_size=(3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Conv2D(filters=128, kernel_size=(3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Flatten(),
    BatchNormalization(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(units=num_classes, activation='softmax')
])

simple_model.summary()

```

Model: "sequential_11"

Layer (type)	Output Shape	
conv2d_31 (Conv2D)	(None, 222, 222, 32)	
max_pooling2d_22 (MaxPooling2D)	(None, 111, 111, 32)	
conv2d_32 (Conv2D)	(None, 109, 109, 64)	
max_pooling2d_23 (MaxPooling2D)	(None, 54, 54, 64)	
conv2d_33 (Conv2D)	(None, 52, 52, 128)	
max_pooling2d_24 (MaxPooling2D)	(None, 26, 26, 128)	
flatten_7 (Flatten)	(None, 86528)	
batch_normalization_30 (BatchNormalization)	(None, 86528)	
dense_35 (Dense)	(None, 512)	
dropout_34 (Dropout)	(None, 512)	
dense_36 (Dense)	(None, 7)	



Total params: 44,745,799 (170.69 MB)

Trainable params: 44,572,743 (170.03 MB)

Non-trainable params: 173,056 (676.00 KB)

In [173...

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=[

EPOCHS = 15
PATIENCE = 7

early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=PATIENCE,
    restore_best_weights=True
)

history1 = model.fit(
    train_gen,
    validation_data=test_gen,
    epochs=EPOCHS,
    callbacks=[early_stopping]
)
```

Epoch 1/15
248/248  **205s** 821ms/step - accuracy: 0.7002 - loss: 0.8100 - val_accuracy: 0.6687 - val_loss: 0.8508

Epoch 2/15
248/248  **211s** 850ms/step - accuracy: 0.7011 - loss: 0.7923 - val_accuracy: 0.6878 - val_loss: 0.7994

Epoch 3/15
248/248  **206s** 832ms/step - accuracy: 0.7006 - loss: 0.7903 - val_accuracy: 0.7038 - val_loss: 0.7943

Epoch 4/15
248/248  **223s** 898ms/step - accuracy: 0.7095 - loss: 0.7680 - val_accuracy: 0.7013 - val_loss: 0.7858

Epoch 5/15
248/248  **205s** 824ms/step - accuracy: 0.7106 - loss: 0.7618 - val_accuracy: 0.6943 - val_loss: 0.8150

Epoch 6/15
248/248  **221s** 892ms/step - accuracy: 0.7088 - loss: 0.7565 - val_accuracy: 0.7169 - val_loss: 0.7493

Epoch 7/15
248/248  **213s** 859ms/step - accuracy: 0.7075 - loss: 0.7484 - val_accuracy: 0.7038 - val_loss: 0.7643

Epoch 8/15
248/248  **205s** 824ms/step - accuracy: 0.7180 - loss: 0.7359 - val_accuracy: 0.7078 - val_loss: 0.7659

Epoch 9/15
248/248  **209s** 844ms/step - accuracy: 0.7184 - loss: 0.7289 - val_accuracy: 0.6958 - val_loss: 0.7634

Epoch 10/15
248/248  **205s** 827ms/step - accuracy: 0.7229 - loss: 0.7358 - val_accuracy: 0.7169 - val_loss: 0.7597


Epoch 11/15
248/248  **216s** 872ms/step - accuracy: 0.7193 - loss: 0.7384 - val_accuracy: 0.6973 - val_loss: 0.8154

Epoch 12/15
248/248  **214s** 861ms/step - accuracy: 0.7222 - loss: 0.7257 - val_accuracy: 0.7149 - val_loss: 0.7635

Epoch 13/15
248/248  **209s** 841ms/step - accuracy: 0.7175 - loss: 0.7342 - val_accuracy: 0.7189 - val_loss: 0.7542

In [174...

```
plt.plot(history1.history['val_loss'],label='validation loss')
plt.plot(history1.history['loss'],label='training loss')
plt.legend()
plt.show()
```



```
In [175... y_pred1 = model.predict(test_gen)
```

63/63 ————— 14s 218ms/step

```
In [176... pred_classes1 = y_pred1.argmax(axis=1)
pred_classes1 = pred_classes1[:test_gen.samples] ## get the first 1992 out of 13
```

```
In [177... ConfusionMatrixDisplay.from_predictions(test_gen.classes, pred_classes1, display
plt.show())
```



```
BatchNormalization(),
MaxPooling2D(pool_size=(2,2)),
Dropout(0.25),

## fully connect all layers
Flatten(),
Dense(512, activation='relu'),
BatchNormalization(),
Dropout(0.5),
Dense(256, activation='relu'),
BatchNormalization(),
Dropout(0.5),
Dense(units=num_classes, activation='softmax')
])

deep_model.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	
conv2d_25 (Conv2D)	(None, 222, 222, 32)	
batch_normalization_22 (BatchNormalization)	(None, 222, 222, 32)	
conv2d_26 (Conv2D)	(None, 220, 220, 32)	
batch_normalization_23 (BatchNormalization)	(None, 220, 220, 32)	
max_pooling2d_19 (MaxPooling2D)	(None, 110, 110, 32)	
dropout_29 (Dropout)	(None, 110, 110, 32)	
conv2d_27 (Conv2D)	(None, 108, 108, 64)	
batch_normalization_24 (BatchNormalization)	(None, 108, 108, 64)	
conv2d_28 (Conv2D)	(None, 106, 106, 64)	
batch_normalization_25 (BatchNormalization)	(None, 106, 106, 64)	
max_pooling2d_20 (MaxPooling2D)	(None, 53, 53, 64)	
dropout_30 (Dropout)	(None, 53, 53, 64)	
conv2d_29 (Conv2D)	(None, 51, 51, 128)	
batch_normalization_26 (BatchNormalization)	(None, 51, 51, 128)	
conv2d_30 (Conv2D)	(None, 49, 49, 128)	
batch_normalization_27 (BatchNormalization)	(None, 49, 49, 128)	
max_pooling2d_21 (MaxPooling2D)	(None, 24, 24, 128)	
dropout_31 (Dropout)	(None, 24, 24, 128)	
flatten_6 (Flatten)	(None, 73728)	
dense_32 (Dense)	(None, 512)	
batch_normalization_28 (BatchNormalization)	(None, 512)	
dropout_32 (Dropout)	(None, 512)	
dense_33 (Dense)	(None, 256)	
batch_normalization_29 (BatchNormalization)	(None, 256)	
dropout_33 (Dropout)	(None, 256)	

dense_34 (Dense)	(None, 7)	
------------------	-----------	--

Total params: 38,174,247 (145.62 MB)

Trainable params: 38,171,815 (145.61 MB)

Non-trainable params: 2,432 (9.50 KB)

```
In [166... deep_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metr


early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

deep_history = deep_model.fit(
    train_gen,
    validation_data=test_gen,
    epochs=20,
    callbacks=[early_stopping]
)
```


Epoch 1/20
248/248  **539s** 2s/step - accuracy: 0.4996 - loss: 1.7182 - val
 _accuracy: 0.1722 - val_loss: 2.3335
 Epoch 2/20
248/248  **535s** 2s/step - accuracy: 0.6324 - loss: 1.0933 - val
 _accuracy: 0.6642 - val_loss: 0.9484
 Epoch 3/20
248/248  **530s** 2s/step - accuracy: 0.6600 - loss: 0.9713 - val
 _accuracy: 0.6662 - val_loss: 0.9227
 Epoch 4/20
248/248  **527s** 2s/step - accuracy: 0.6713 - loss: 0.9210 - val
 _accuracy: 0.6320 - val_loss: 0.9388
 Epoch 5/20
248/248  **522s** 2s/step - accuracy: 0.6876 - loss: 0.8624 - val
 _accuracy: 0.6732 - val_loss: 0.8596
 Epoch 6/20
248/248  **534s** 2s/step - accuracy: 0.6918 - loss: 0.8379 - val
 _accuracy: 0.4649 - val_loss: 1.3769
 Epoch 7/20
248/248  **524s** 2s/step - accuracy: 0.6983 - loss: 0.8189 - val
 _accuracy: 0.7083 - val_loss: 0.7840
 Epoch 8/20
248/248  **526s** 2s/step - accuracy: 0.7077 - loss: 0.7978 - val
 _accuracy: 0.7003 - val_loss: 0.8266
 Epoch 9/20
248/248  **528s** 2s/step - accuracy: 0.7077 - loss: 0.7834 - val
 _accuracy: 0.6355 - val_loss: 0.9728
 Epoch 10/20
248/248  **524s** 2s/step - accuracy: 0.6778 - loss: 0.9133 - val
 _accuracy: 0.6762 - val_loss: 0.9710
 Epoch 11/20
248/248  **524s** 2s/step - accuracy: 0.6890 - loss: 0.8515 - val
 _accuracy: 0.6928 - val_loss: 0.8562
 Epoch 12/20
248/248  **528s** 2s/step - accuracy: 0.7037 - loss: 0.7985 - val
 _accuracy: 0.7028 - val_loss: 0.8368
 Epoch 13/20
248/248  **527s** 2s/step - accuracy: 0.7058 - loss: 0.7855 - val
 _accuracy: 0.6526 - val_loss: 0.8786
 Epoch 14/20
248/248  **533s** 2s/step - accuracy: 0.7127 - loss: 0.7733 - val
 _accuracy: 0.6109 - val_loss: 0.9696
 Epoch 15/20
248/248  **525s** 2s/step - accuracy: 0.7101 - loss: 0.7679 - val
 _accuracy: 0.7174 - val_loss: 0.8005
 Epoch 16/20
248/248  **519s** 2s/step - accuracy: 0.7167 - loss: 0.7519 - val
 _accuracy: 0.7078 - val_loss: 0.7908
 Epoch 17/20
248/248  **531s** 2s/step - accuracy: 0.7246 - loss: 0.7383 - val
 _accuracy: 0.6857 - val_loss: 0.8940

In [167...

```
plt.plot(deep_history.history['val_loss'],label='validation loss')
plt.plot(deep_history.history['loss'],label='training loss')
plt.legend()
plt.show()
```




```
In [168... y_pred_deep = model.predict(test_gen)
```

63/63 ————— 16s 239ms/step

```
In [169... pred_classes_deep = y_pred1.argmax(axis=1)
pred_classes_deep = pred_classes1[:test_gen.samples] ## get the first 1992 out of
```

```
In [170... ConfusionMatrixDisplay.from_predictions(test_gen.classes, pred_classes_deep, display=
plt.show())
```



```
In [171... t_lost, t_acc = deep_model.evaluate(test_gen, verbose=1)
t_acc
```

63/63 ————— 23s 359ms/step - accuracy: 0.7083 - loss: 0.7840

```
Out[171... 0.7083333134651184
```

Building the Transfer Learning Models

```
In [155... def create_transfer_learning_model(model_name):

    if model_name == 'MobileNetV2':
        base_model = MobileNetV2(weights = 'imagenet', include_top=False, input_
    elif model_name == 'ResNet50':
        base_model = ResNet50(weights = 'imagenet', include_top=False, input_sha
    elif model_name == 'VGG16':
        base_model = VGG16(weights = 'imagenet', include_top=False, input_shape=
    else:
        print('Unknown Transfer Learning model')

    ## build the model
    base_model.trainable = False

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dense(512,activation='relu'),
        BatchNormalization(),
        Dropout(0.5),
        Dense(256,activation='relu'),
```

```

        # BatchNormalization(),
        Dropout(0.5),
        Dense(512,activation='relu'),
        BatchNormalization(),
        Dropout(0.5),
        # Dense(256,activation='relu'),
        # BatchNormalization(),
        # Dropout(0.5),
        Dense(units=num_classes,activation='softmax')
    ])

    return model

# Create transfer learning model with MobileNetV2
mobilenet_model = create_transfer_learning_model('MobileNetV2')
mobilenet_model.summary()

```

Model: "sequential_8"

Layer (type)	Output Shape	
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	
global_average_pooling2d_3 (GlobalAveragePooling2D)	(None, 1280)	
dense_25 (Dense)	(None, 512)	
batch_normalization_12 (BatchNormalization)	(None, 512)	
dropout_21 (Dropout)	(None, 512)	
dense_26 (Dense)	(None, 256)	
dropout_22 (Dropout)	(None, 256)	
dense_27 (Dense)	(None, 512)	
batch_normalization_13 (BatchNormalization)	(None, 512)	
dropout_23 (Dropout)	(None, 512)	
dense_28 (Dense)	(None, 7)	



Total params: 3,184,455 (12.15 MB)
 Trainable params: 924,423 (3.53 MB)
 Non-trainable params: 2,260,032 (8.62 MB)

```

In [49]: mobilenet_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy'

early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

```

```
history = mobilenet_model.fit(  
    train_gen,  
    validation_data=test_gen,  
    epochs=20,  
    callbacks=[early_stopping]  
)
```

Epoch 1/20
248/248  165s 648ms/step - accuracy: 0.5294 - loss: 1.6506 - val_accuracy: 0.6993 - val_loss: 0.8835

Epoch 2/20
248/248  156s 630ms/step - accuracy: 0.6508 - loss: 1.1150 - val_accuracy: 0.7134 - val_loss: 0.8346

Epoch 3/20
248/248  156s 630ms/step - accuracy: 0.6783 - loss: 0.9654 - val_accuracy: 0.7103 - val_loss: 0.7904

Epoch 4/20
248/248  157s 631ms/step - accuracy: 0.6901 - loss: 0.9048 - val_accuracy: 0.7314 - val_loss: 0.7826

Epoch 5/20
248/248  157s 634ms/step - accuracy: 0.6969 - loss: 0.8694 - val_accuracy: 0.7209 - val_loss: 0.8054

Epoch 6/20
248/248  153s 617ms/step - accuracy: 0.6961 - loss: 0.8478 - val_accuracy: 0.7159 - val_loss: 0.7806

Epoch 7/20
248/248  159s 639ms/step - accuracy: 0.7094 - loss: 0.8274 - val_accuracy: 0.7108 - val_loss: 0.7898

Epoch 8/20
248/248  207s 836ms/step - accuracy: 0.7135 - loss: 0.8013 - val_accuracy: 0.7294 - val_loss: 0.7728

Epoch 9/20
248/248  226s 911ms/step - accuracy: 0.7094 - loss: 0.8063 - val_accuracy: 0.7078 - val_loss: 0.7990

Epoch 10/20
248/248  216s 871ms/step - accuracy: 0.7172 - loss: 0.7884 - val_accuracy: 0.7214 - val_loss: 0.7744

Epoch 11/20
248/248  207s 835ms/step - accuracy: 0.7179 - loss: 0.7967 - val_accuracy: 0.7344 - val_loss: 0.7604

Epoch 12/20
248/248  208s 840ms/step - accuracy: 0.7242 - loss: 0.7811 - val_accuracy: 0.7304 - val_loss: 0.7595

Epoch 13/20
248/248  200s 807ms/step - accuracy: 0.7268 - loss: 0.7620 - val_accuracy: 0.7369 - val_loss: 0.7575

Epoch 14/20
248/248  206s 830ms/step - accuracy: 0.7302 - loss: 0.7647 - val_accuracy: 0.7088 - val_loss: 0.7655

Epoch 15/20
248/248  208s 838ms/step - accuracy: 0.7283 - loss: 0.7512 - val_accuracy: 0.7420 - val_loss: 0.7258

Epoch 16/20
248/248  201s 812ms/step - accuracy: 0.7291 - loss: 0.7476 - val_accuracy: 0.7475 - val_loss: 0.7244

Epoch 17/20
248/248  209s 841ms/step - accuracy: 0.7264 - loss: 0.7548 - val_accuracy: 0.7239 - val_loss: 0.7634


Epoch 18/20
248/248  211s 850ms/step - accuracy: 0.7338 - loss: 0.7405 - val_accuracy: 0.7500 - val_loss: 0.7342

Epoch 19/20
248/248  215s 865ms/step - accuracy: 0.7331 - loss: 0.7422 - val_accuracy: 0.7359 - val_loss: 0.7475


Epoch 20/20
248/248  207s 835ms/step - accuracy: 0.7344 - loss: 0.7371 - val_accuracy: 0.7380 - val_loss: 0.7518

```
In [50]: pd.DataFrame(history.history).plot()
```


Out[50]: <Axes: >



```
In [52]: plt.plot(history.history['val_loss'],label='validation loss')
plt.plot(history.history['loss'],label='training loss')
plt.legend()
plt.show()
```



```
In [181... plt.plot(history.history['accuracy'],label='Training Accuracy')
plt.plot(history.history['val_accuracy'],label='Validation Accuracy')
plt.legend()
plt.show()
```




```
In [53]: y_pred = mobilenet_model.predict(test_gen)
```

63/63 ————— 24s 357ms/step

```
In [54]: pred_classes = y_pred.argmax(axis=1)
pred_classes = pred_classes[:test_gen.samples] ## get the first 1992 out of 1394
```

```
In [55]: ConfusionMatrixDisplay.from_predictions(test_gen.classes,pred_classes, display_1
plt.show())
```

```
In [56]: t_lost,t_acc = mobilenet_model.evaluate(test_gen, verbose=1)
         t_acc
```

63/63 ————— 21s 333ms/step - accuracy: 0.7475 - loss: 0.7244

```
Out[56]: 0.7474899888038635
```

MobileNetV2 model has the highest accuracy among the 3 models.

Testing the Best CNN Model's Predictions with Sample Images

```
In [195... sample_images = [f for f in os.listdir(SAMPLE_DIR) if f.endswith(('.jpg', '.jpeg'))

fig, axes = plt.subplots(2,4,figsize=(16,8))
axes = axes.flatten()

for i, img_name in enumerate(sample_images[:8]):
    print(img_name)
    img_path = os.path.join(SAMPLE_DIR, img_name)
    img = image.load_img(img_path, target_size=(IMG_HEIGHT, IMG_WIDTH))
    img_array = image.img_to_array(img)
    img_array = np.expand_dims(img_array, axis=0)
    img_array = img_array/255.0


    pred = loaded_model.predict(img_array, verbose=0)[0]
    pred_class_idx = np.argmax(pred)
    confidence = pred[pred_class_idx] * 100
    pred_label = class_names[pred_class_idx]
    # print(pred)
```

```

axes[i].imshow(img)
axes[i].set_title(f'{pred_label}\n({confidence:.1f}%)confident')
axes[i].axis('off')
plt.suptitle('Model Prediction on sample images')
plt.tight_layout()
plt.show()

```

ISIC_0029325.jpg
 ISIC_0029334.jpg
 ISIC_0029341.jpg
 ISIC_0029371.jpg
 ISIC_0029377.jpg
 ISIC_0029382.jpg
 ISIC_0029386.jpg
 ISIC_0029396.jpg



In [59]: `mobilenet_model.save('./good_model_skin_cancer.keras')`

In [144... `patient_df.to_csv('patient_data_cleaned.csv')`

In [179... `loaded_model = tf.keras.models.load_model('good_model_skin_cancer.keras')`

In [180... `t_lost,t_acc = loaded_model.evaluate(test_gen, verbose=1)`
`t_acc`

63/63 ————— 25s 366ms/step - accuracy: 0.7475 - loss: 0.7244

Out[180... 0.7474899888038635